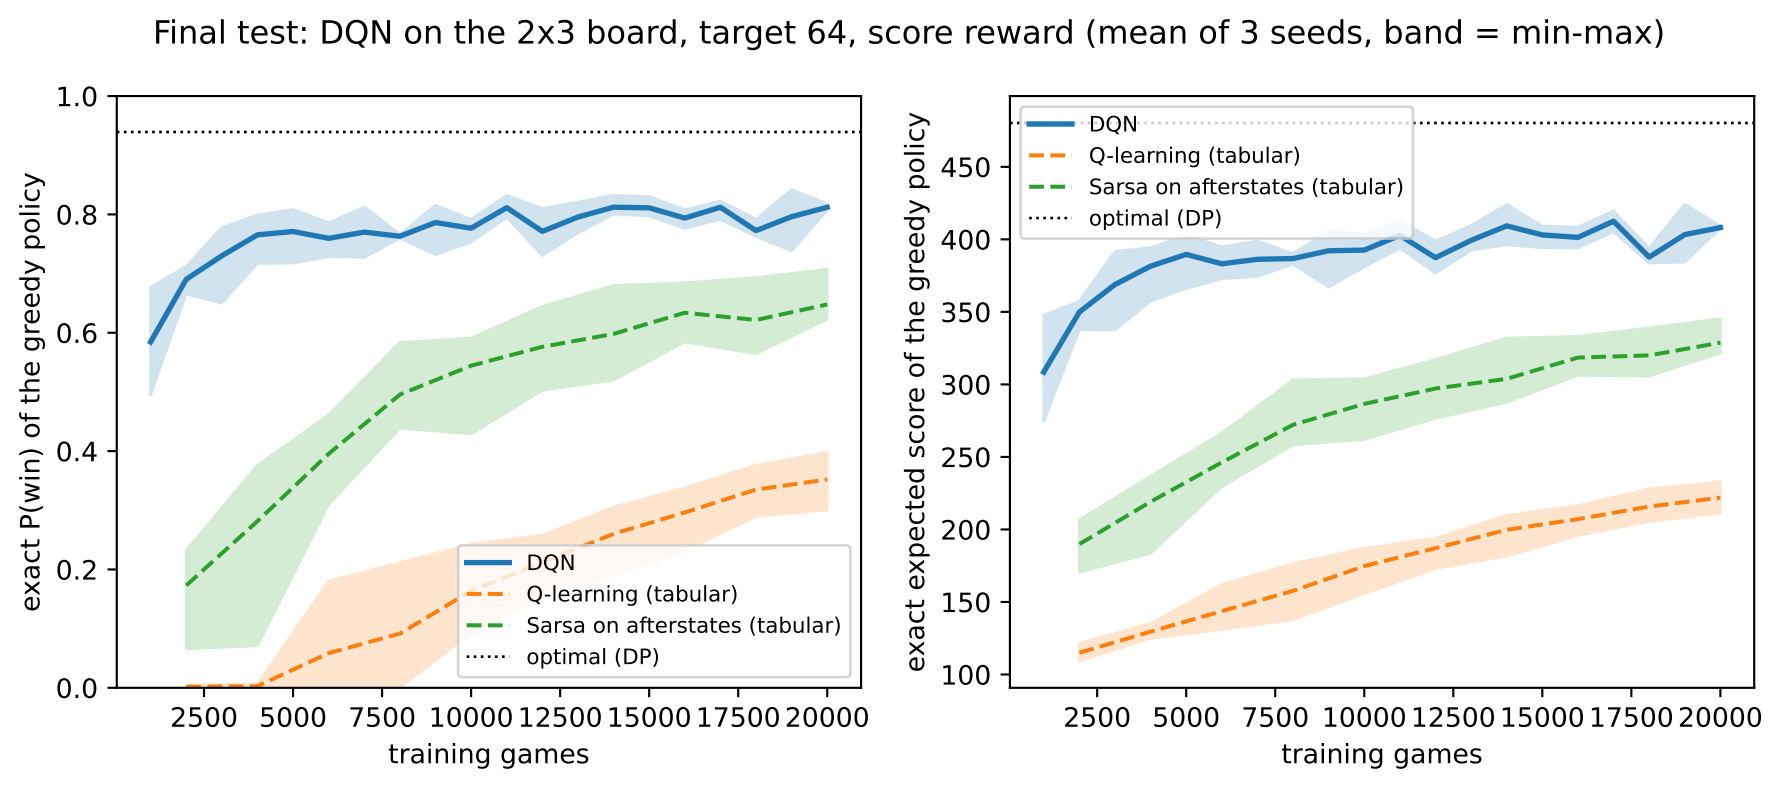
Final test: DQN on the 2x3 board, target 64, score reward (mean of 3 seeds, band = min-max)
exact P(win) of the greedy policy
1.0
0.8
0.6
0.4
0.2
0.0
2500
5000
7500
10000
12500
15000
17500
20000
training games
DQN
Q-learning (tabular)
Sarsa on afterstates (tabular)
optimal (DP)
exact expected score of the greedy policy
450
400
350
300
250
200
150
100
2500
5000
7500
10000
12500
15000
17500
20000
training games
DQN
Q-learning (tabular)
Sarsa on afterstates (tabular)
optimal (DP)