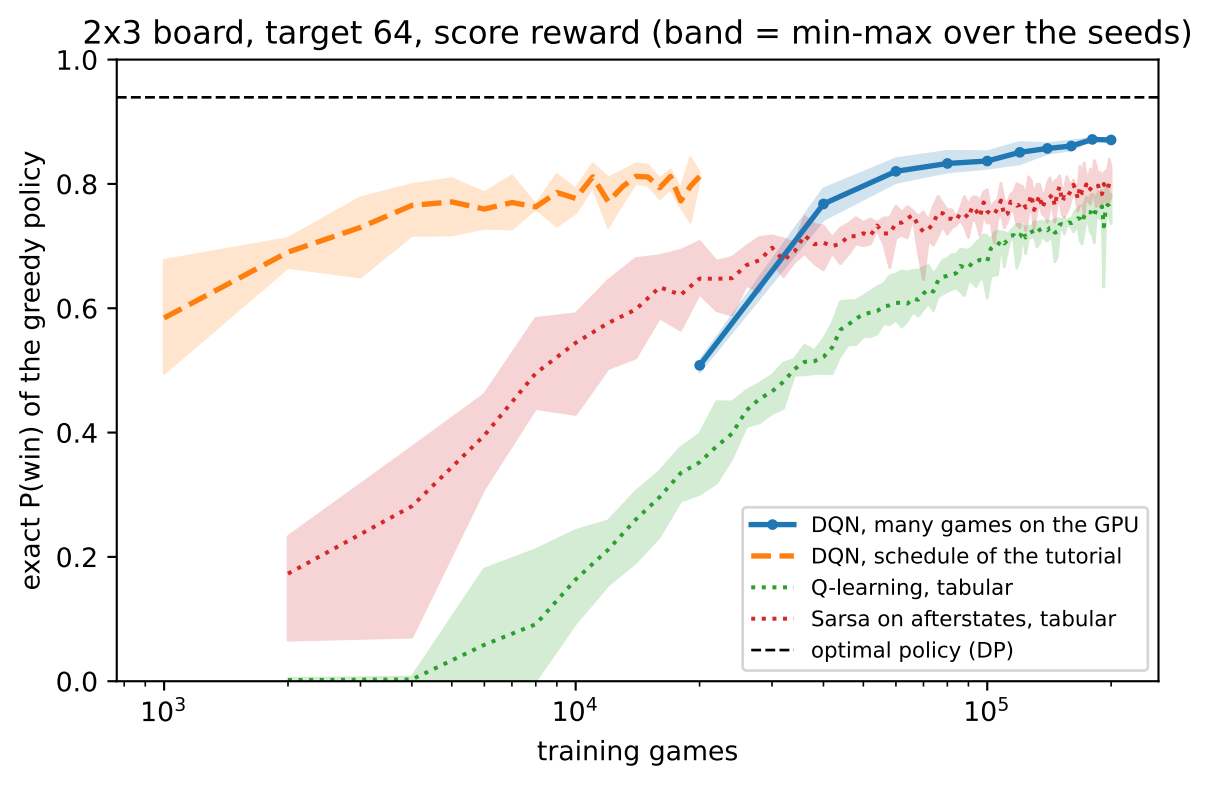
2x3 board, target 64, score reward (band = min-max over the seeds)
exact P(win) of the greedy policy
training games
1.0
0.8
0.6
0.4
0.2
0.0
10^3
10^4
10^5
DQN, many games on the GPU
DQN, schedule of the tutorial
Q-learning, tabular
Sarsa on afterstates, tabular
optimal policy (DP)