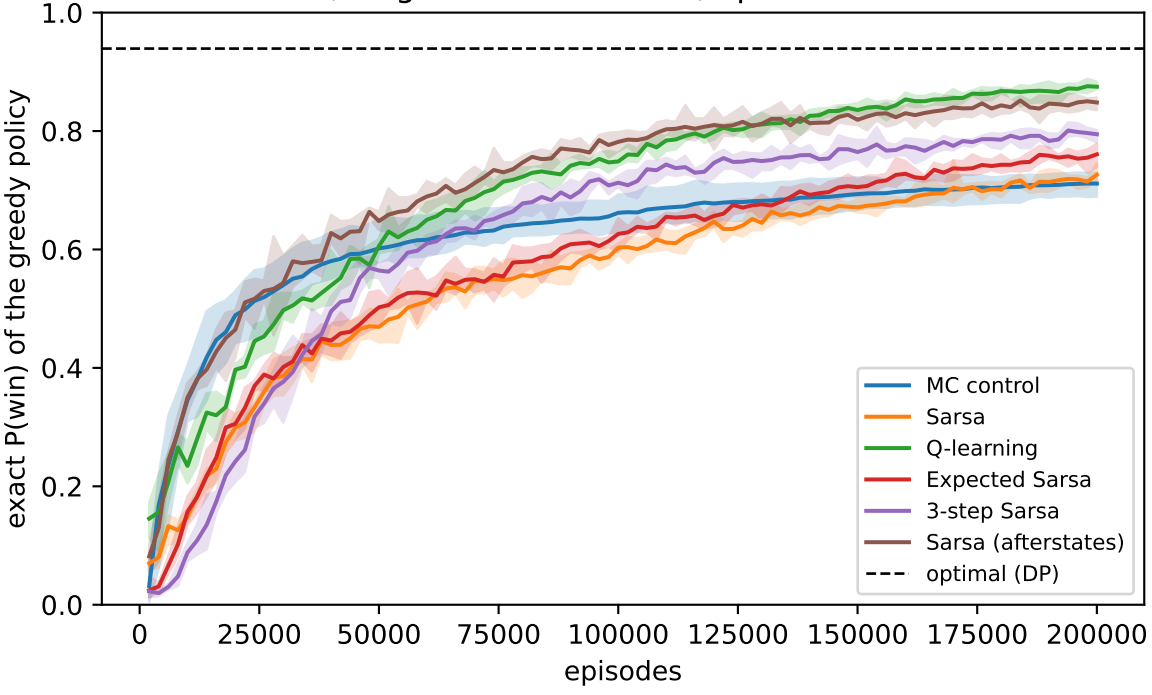
2x3 board, target 64: win reward, optimistic initialisation
exact P(win) of the greedy policy
episodes
1.0
0.8
0.6
0.4
0.2
0.0
0
25000
50000
75000
100000
125000
150000
175000
200000
MC control
Sarsa
Q-learning
Expected Sarsa
3-step Sarsa
Sarsa (afterstates)
optimal (DP)