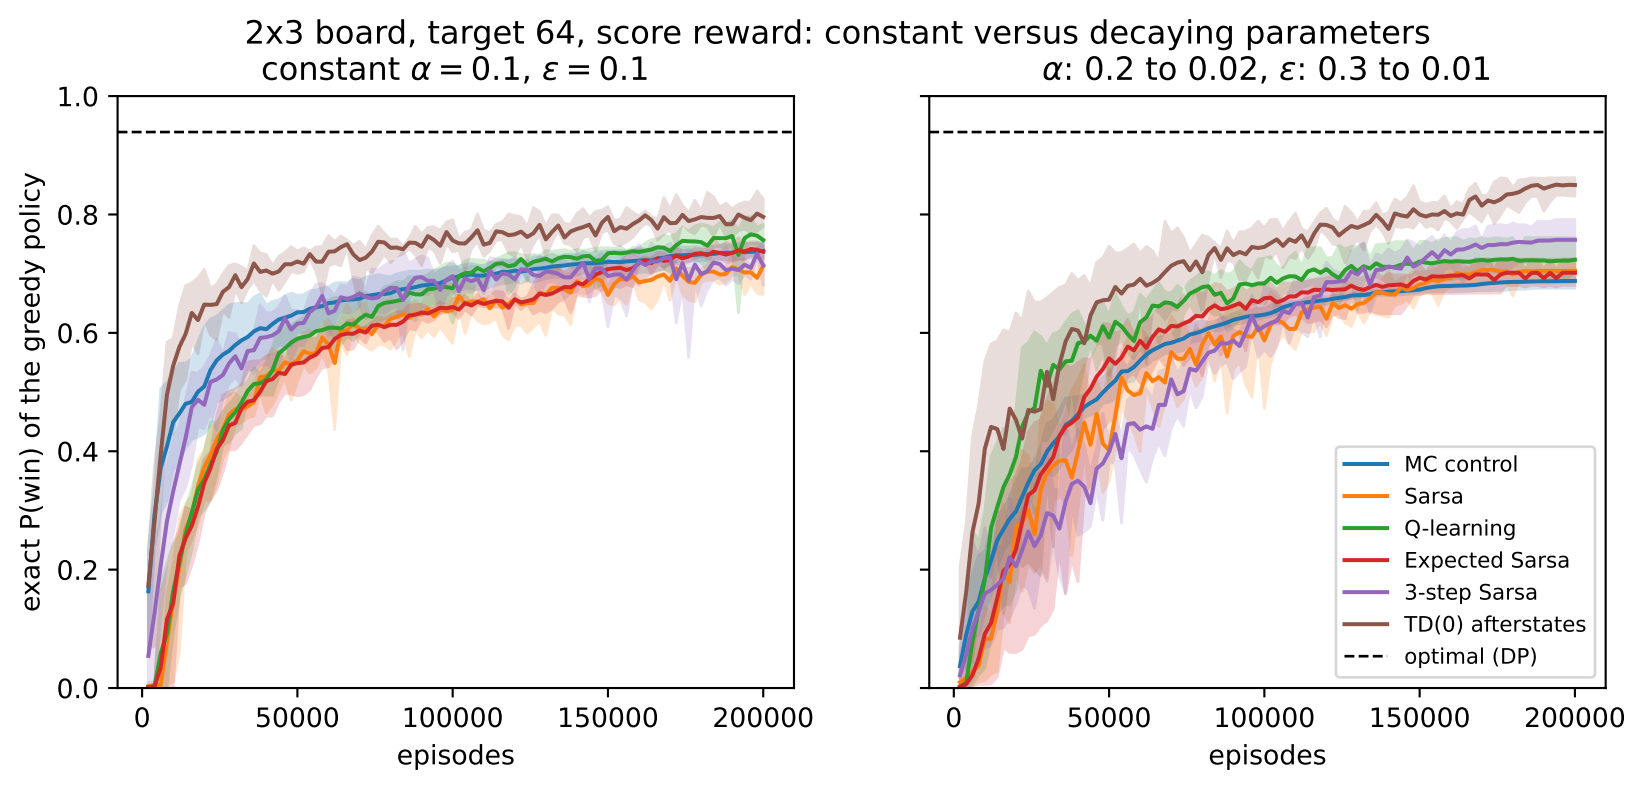
2x3 board, target 64, score reward: constant versus decaying parameters
constant $\alpha = 0.1$, $\varepsilon = 0.1$
α: 0.2 to 0.02, ε: 0.3 to 0.01
exact P(win) of the greedy policy
1.0
0.8
0.6
0.4
0.2
0.0
0
50000
100000
150000
200000
episodes
MC control
Sarsa
Q-learning
Expected Sarsa
3-step Sarsa
TD(0) afterstates
optimal (DP)